

# **DISTRIBUTED COMPUTING FOR NEW BLOODS**

**RAYMOND TAY**

**I WORK IN HEWLETT-PACKARD LABS SINGAPORE**

**@RAYMONDTAYBL**



# **What is a Distributed System?**



# What is a Distributed System?

*It's been  
there for  
a long  
time*



# What is a Distributed System?

*It's been  
there for  
a long  
time*

**You did not realize it**



# **DISTRIBUTED SYSTEMS ARE THE NEW FRONTIER**

**whether YOU like it or not**



# **DISTRIBUTED SYSTEMS ARE THE NEW FRONTIER**

**LAN Games**

**Mobile**

**Databases**

**ATMs**

**Social Media**

**whether YOU like it or not**



# **WHAT IS THE ESSENCE OF DISTRIBUTED SYSTEMS?**



**WHAT IS THE  
ESSENCE OF  
DISTRIBUTED  
SYSTEMS?**

**IT IS ABOUT THE  
ATTEMPT TO  
OVERCOME**

- \* **INFORMATION  
TRAVEL**
- \* **WHEN  
INDEPENDENT  
PROCESSES FAIL**




# Information flows at speed of light !





When  
independent  
things **DONT** fail.

I've sent it.



I've received it.




I've received it.

I've sent it.




**When independent  
things fail,  
independently.**



**NETWORK FAILURE**



**When independent  
things fail,  
independently.**



**NODE FAILURE**



**There is no difference  
between a slow node  
and a dead node**



# Why do we NEED it?

## **Scalability**

**when the needs of the  
system outgrows what a  
single node can provide**



# Why do we NEED it?

## **Scalability**

**when the needs of the  
system outgrows what a  
single node can provide**

## **Availability**

**Enabling resilience  
when a node fails**



**IS IT REALLY THAT  
HARD?**



# IS IT REALLY THAT HARD?

The answer lies in  
knowing what is  
possible and what is not  
possible...



**IS IT REALLY THAT  
HARD?**

**Even widely accepted  
facts are challenged...**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- ✱ The network is reliable

## AWS EBS Outage

On April 21, 2011, AWS suffered unavailability for 12 hours,<sup>2</sup> causing hundreds of high-profile Web sites to go offline. As a part of normal AWS scaling activities, Amazon engineers had shifted traffic away from a router in the EBS (Elastic Block Store) network in a single U.S. East AZ (Availability Zone), but, due to incorrect routing policies:

...many EBS nodes in the affected Availability Zone were completely isolated from other EBS nodes in its cluster. Unlike a normal network interruption, this change disconnected both the primary and secondary network simultaneously, leaving the affected nodes completely isolated from one another.



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**




# THE FALLACIES OF DISTRIBUTED COMPUTING



## Cyber security budget allocation

Figure 16 summarizes six layers in a typical multi-layered IT security infrastructure for all benchmarked companies. Each bar reflects the percentage dedicated spending according to the presented layer. The network layer receives the highest allocation at 40 percent of total dedicated IT security funding. At only five percent, the host layer receives the lowest funding level. The percentage allocations to physical layer activities is highest for critical infrastructure companies such as communications, energy and utilities and lowest for retail, hospitality and consumer product companies.

**Figure 16. Budgeted or earmarked spending according to six IT security layers**



**2013 cost of cyber  
crime study:  
United States**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero

**Ingo Rammer on  
latency vs  
bandwidth**

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**

*"But I think that it's really interesting to see that the end-to-end bandwidth increased by 1468 times within the last 11 years while the latency (the time a single ping takes) has only been improved tenfold. If this wouldn't be enough, there is even a natural cap on latency. The minimum round-trip time between two points of this earth is determined by the maximum speed of information transmission: the speed of light. At roughly 300,000 kilometers per second ( $3.6 * 10^{12}$  teraangstrom per fortnight), it will always take at least 30 milliseconds to send a ping from Europe to the US and back, even if the processing would be done in real time."*



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero
- \* Bandwidth is infinite

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero
- \* Bandwidth is infinite
- \* Topology doesn't change

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero
- \* Bandwidth is infinite
- \* Topology doesn't change
- \* There is one administrator

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero
- \* Bandwidth is infinite
- \* Topology doesn't change
- \* There is one administrator
- \* Transport cost is zero

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# THE FALLACIES OF DISTRIBUTED COMPUTING

- \* The network is reliable
- \* The network is secure
- \* The network is homogeneous
- \* Latency is zero
- \* Bandwidth is infinite
- \* Topology doesn't change
- \* There is one administrator
- \* Transport cost is zero

**Peter Deutsch and  
other fellows  
at Sun  
Microsystems**



# WHAT WAS TRIED IN THE PAST?

- \* **DISTRIBUTED OBJECTS**
- \* **REMOTE PROCEDURE CALL**
- \* **DISTRIBUTED SHARED MUTABLE STATE**



# WHAT WAS ATTEMPTED?

- \* **DISTRIBUTED OBJECTS**
- \* **REMOTE PROCEDURE CALL**
- \* **DISTRIBUTED SHARED MUTABLE STATE**
  - Typically, programming constructs to “abstract” the fact that there are local and distributed objects
  - Ignores latencies
  - Handles failures with this attitude → “i dunno what to do...YOU (i.e. invoker) handle it”.



# WHAT WAS ATTEMPTED?

- \* **DISTRIBUTED OBJECTS**
- \* **REMOTE PROCEDURE CALL**
- \* **DISTRIBUTED SHARED MUTABLE STATE**
- **Assumes the synchronous processing model**
- **Asynchronous RPC tries to model after synchronous RPC...**



# WHAT WAS ATTEMPTED?

- \* DISTRIBUTED OBJECTS
- \* REMOTE PROCEDURE CALL
- \* DISTRIBUTED SHARED MUTABLE STATE
  - Found in Distributed Shared Memory Systems i.e.  
“1” address space partitioned into “x” address spaces for “x” nodes
  - e.g. JavaSpaces ⇒ Danger of 2 independent processes to commit successfully.



YOU CAN APPRECIATE  
THAT IT IS  
VERY HARD TO SOLVE  
IN ITS ENTIRETY



The question really  
is “How can we do  
better?”



The question really  
is “How can we do  
better?”

To start, we  
need to  
understand 2  
results



# FLP IMPOSSIBILITY RESULT

The FLP result shows that in an asynchronous model, where one processor might crash, there is no distributed algorithm to solve the consensus problem.



**Consensus is a fundamental problem in fault-tolerant distributed systems. Consensus involves multiple servers agreeing on values. Once they reach a decision on a value, that decision is final.**


**Typical consensus algorithms make progress when any majority of their servers are available; for example, a cluster of 5 servers can continue to operate even if 2 servers fail. If more servers fail, they stop making progress (but will never return an incorrect result).**

**Paxos (1989), Raft (2013)**



# CAP THEOREM


**CAP CONJECTURE**  
(2000) established  
as a theorem in 2002





# CAP THEOREM

**CAP CONJECTURE**  
(2000) established  
as a theorem in 2002




- **Consistency  $\Rightarrow$  all nodes should see the same data, eventually**



# CAP THEOREM

**CAP CONJECTURE**  
(2000) established  
as a theorem in 2002




- **Consistency  $\Rightarrow$  all nodes should see the same data, eventually**
- **Availability  $\Rightarrow$  System is still working when node(s) fails**



# CAP THEOREM

**CAP CONJECTURE**  
(2000) established  
as a theorem in 2002




- **Consistency  $\Rightarrow$  all nodes should see the same data, eventually**
- **Availability  $\Rightarrow$  System is still working when node(s) fails**



# CAP THEOREM

**CAP CONJECTURE**  
(2000) established  
as a theorem in 2002



- **Consistency**  $\Rightarrow$  all nodes should see the same data, eventually
- **Availability**  $\Rightarrow$  System is still working when node(s) fails
- **Partition-Tolerance**  $\Rightarrow$  System is still working on arbitrary message loss



**How does knowing  
FLP and CAP help me?**



# How does knowing FLP and CAP help me?

It's really about asking which  
would you sacrifice when a  
system fails? Consistency  
or Availability?



You can choose C over A

It needs to preserve reads  
and writes i.e. **linearizability**

i.e. A **read** will return the last  
completed write (on ANY  
replica) e.g. **Two-Phase-Commit**



**You can choose A over C**

**It might return stale  
reads ...**



You can choose A over C

It might return stale  
reads ...

DynamoDB uses **vector clocks**;  
Cassandra uses a **clever** form  
of **last-write-wins**



You cannot **NOT** choose P

Partition Tolerance is **mandatory**  
in distributed systems



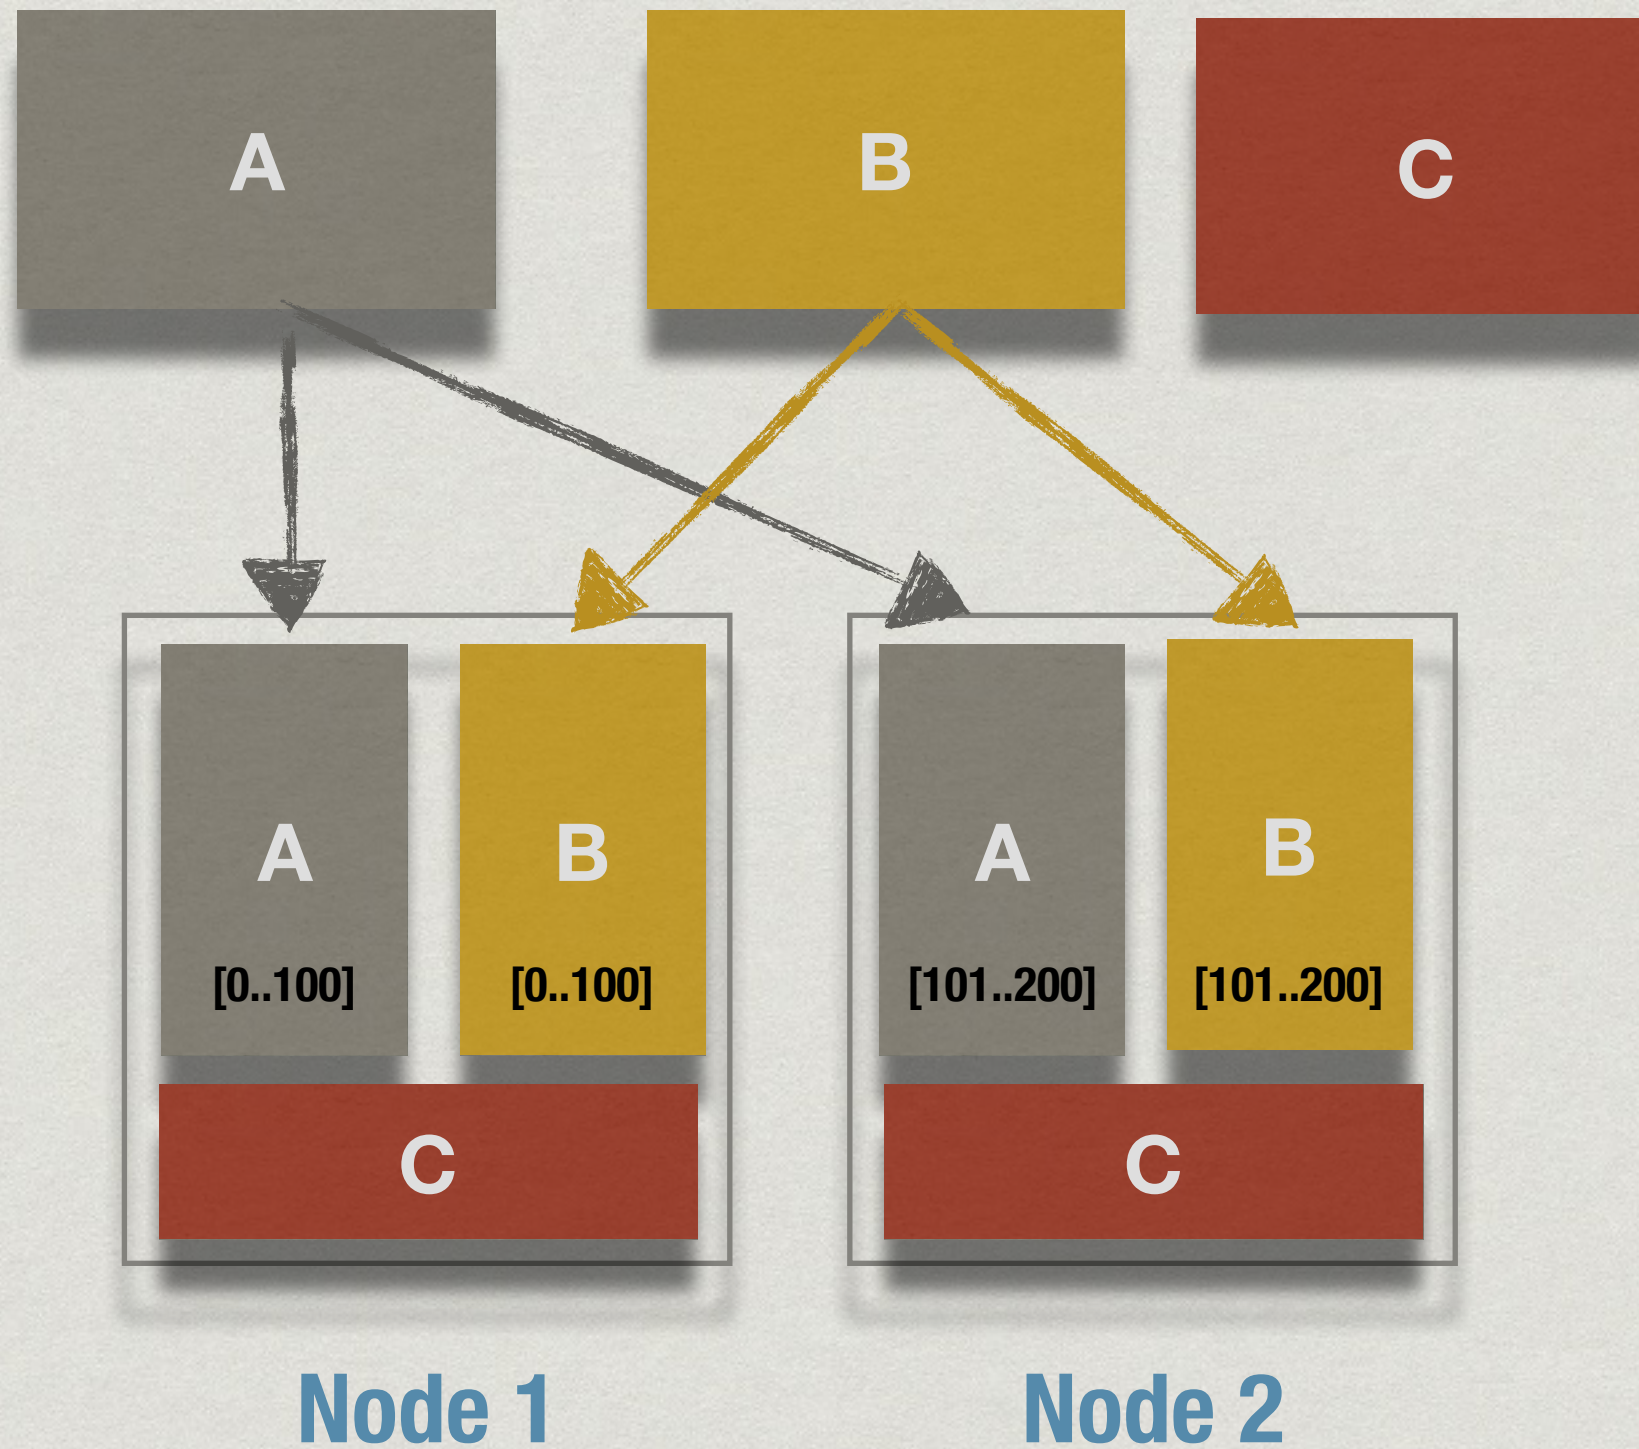
**To scale, partition**

**To resilient, replicate**



## Partitioning

## Replication





**Replication is strongly**  
**related to fault-tolerance**

**Fault-tolerance relies**  
**on reaching consensus**

**Paxos (1989), ZAB, Raft**  
**(2013)**



# Consistent Hashing is used often PARTITIONING and REPLICATION

C-H commonly used in load-  
balancing web objects. In  
DynamoDB Cassandra, Riak and  
Memcached



- **If your problem can fit into memory of a single machine, you don't need a distributed system**
- **If you really need to design / build a distributed system, the following helps:**
  - **Design for failure**
  - **Use the FLP and CAP to critique systems**
  - **Algorithms that work for single-node may not work in distributed mode**
  - **Avoid coordination of nodes**
- **Learn to estimate your capacity**



# Jeff Dean - Google

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns



# REFERENCES

- <http://queue.acm.org/detail.cfm?id=2655736> “The network is reliable”
- <http://cacm.acm.org/blogs/blog-cacm/83396-errors-in-database-systems-eventual-consistency-and-the-cap-theorem/fulltext>
- [http://mvdirona.com/jrh/talksAndPapers/JamesRH\\_Lisa.pdf](http://mvdirona.com/jrh/talksAndPapers/JamesRH_Lisa.pdf)
- “Towards Robust Distributed Systems,” <http://www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf>
- “Why Do Computers Stop and What Can be Done About It,” Tandem Computers Technical Report 85.7, Cupertino, Ca., 1985. <http://www.hpl.hp.com/techreports/tandem/TR-85.7.pdf>
- <http://www.cs.cornell.edu/projects/ladis2009/talks/dean-keynote-ladis2009.pdf>
- <http://codahale.com/you-cant-sacrifice-partition-tolerance/>
- <http://dl.acm.org/citation.cfm?id=258660> “Consistent hashing and random trees: distributed caching protocols for relieving hot spots on the World Wide Web”